


Big Data Summer School

Graph Systems

Research Goal

**Compute on graphs with billions of edges,
in *a reasonable time*, on a single PC.**


- *Reasonable* = close to numbers previously reported for distributed systems in the literature.



Experiment PC: Mac Mini (2012)

Pagerank

- Link-based analysis



Pagerank

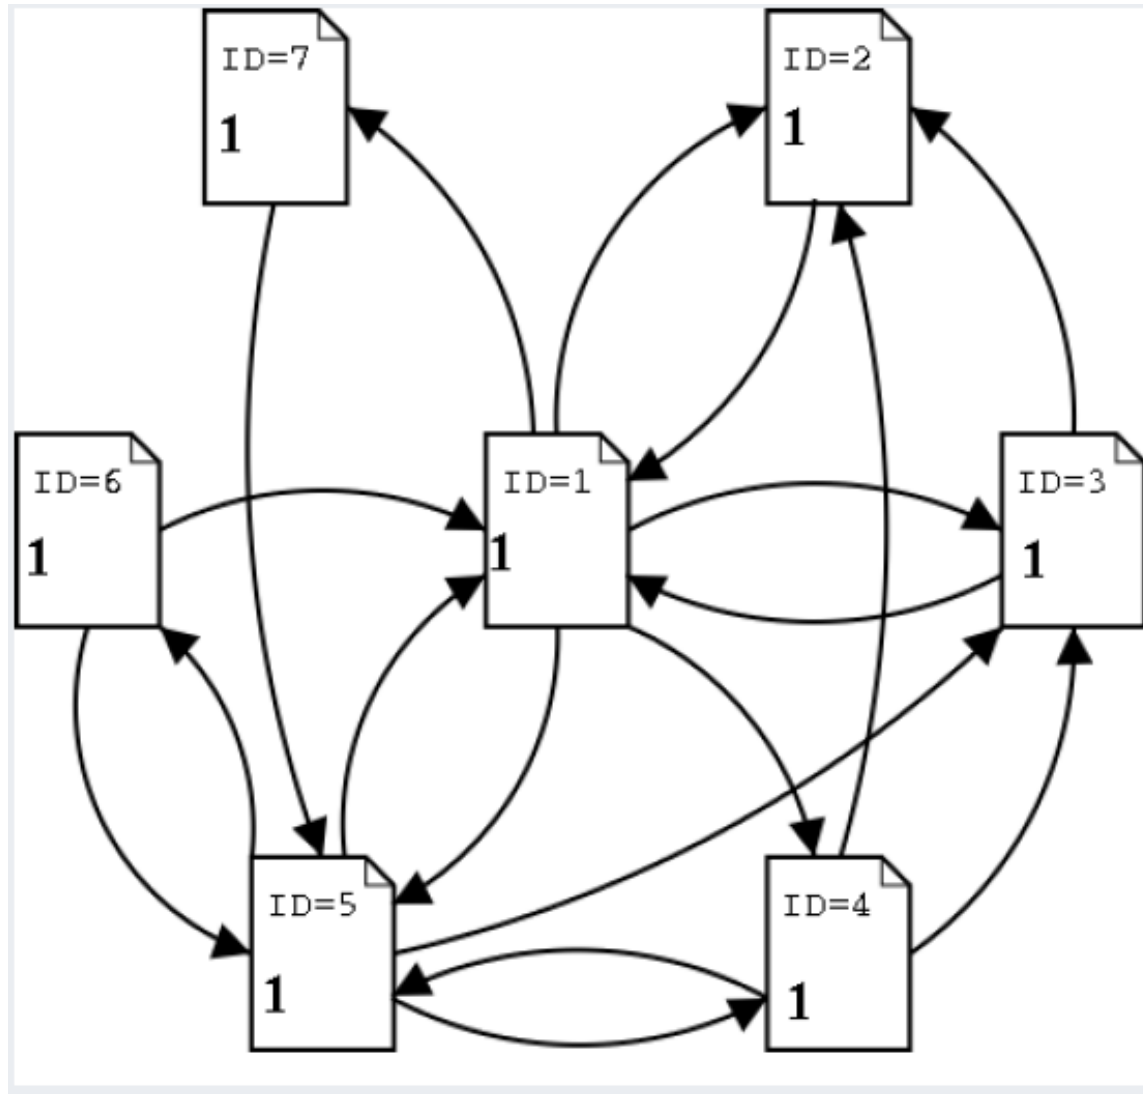
- Link-based analysis

$$PR(u) = \sum_{v \in B_u} \frac{PR(v)}{L(v)}$$

- B_u : the set containing all pages linking to page u
- $L(v)$: the number of links from page v .

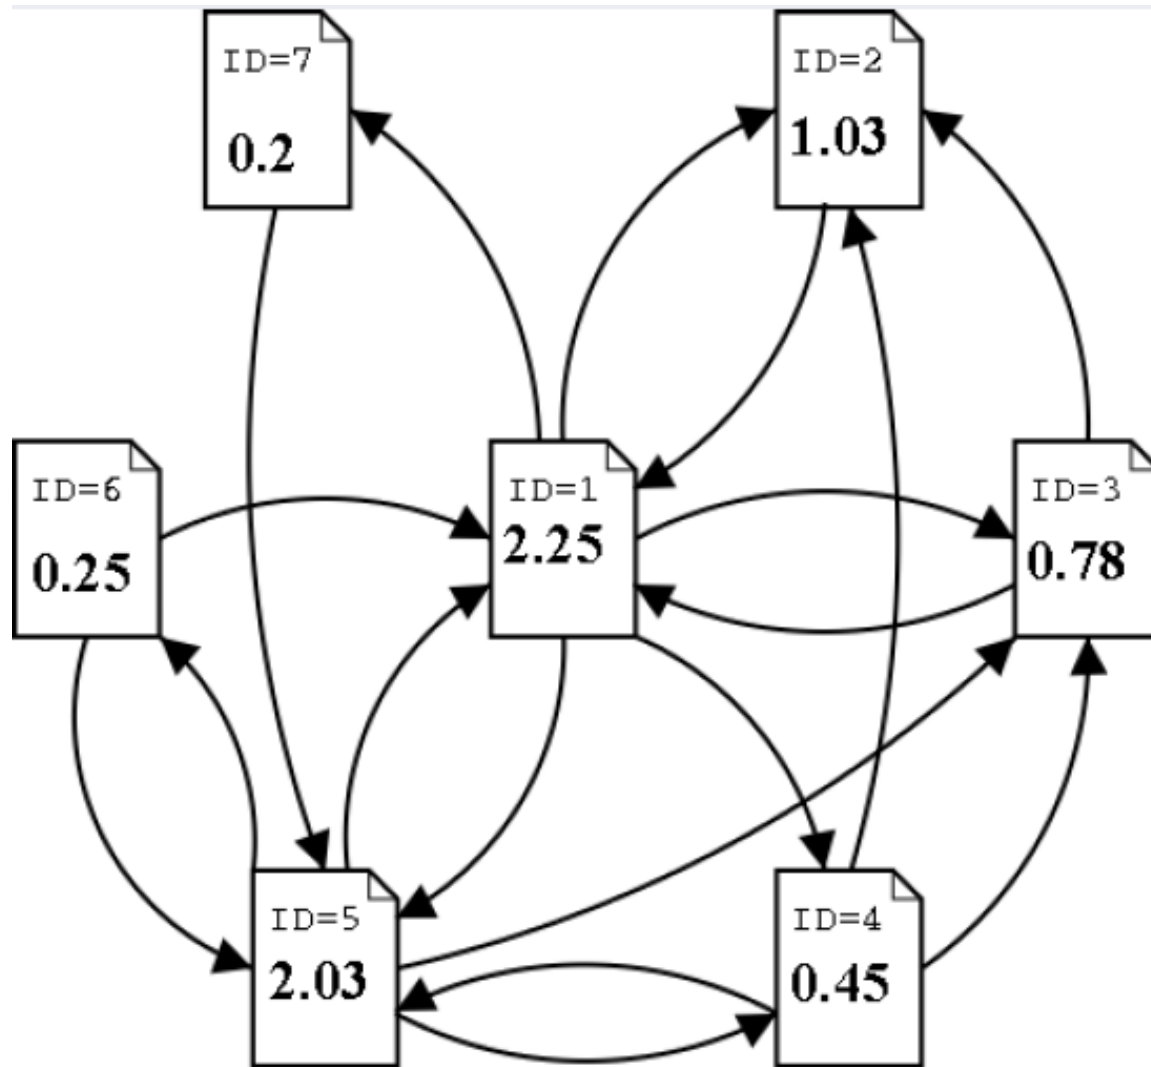
$$PR(u) = \sum_{v \in B_u} \frac{PR(v)}{L(v)}$$

Pagerank



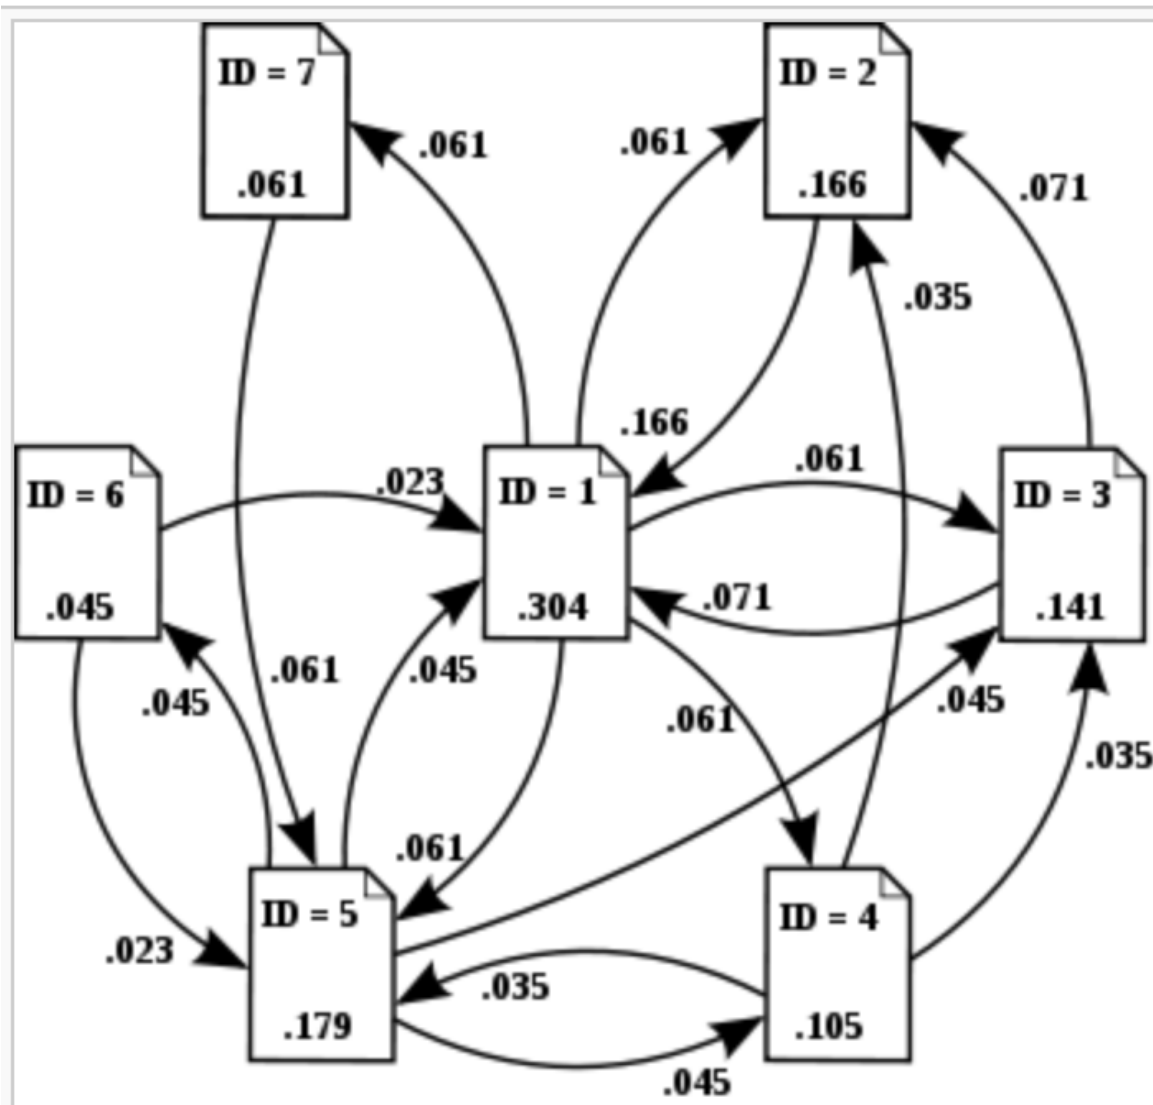
$$PR(u) = \sum_{v \in B_u} \frac{PR(v)}{L(v)}$$

Pagerank



$$PR(u) = \sum_{v \in B_u} \frac{PR(v)}{L(v)}$$


Pagerank




PageRank

- Iterations
 - Gather
 - Apply
 - Scatter


Example: SSSP – Dijkstra's Algorithm




Example: SSSP – Dijkstra's Algorithm




Example: SSSP – Dijkstra's Algorithm




Example: SSSP – Dijkstra's Algorithm



Example: SSSP – Dijkstra's Algorithm




Example: SSSP – Dijkstra's Algorithm




Computational Model

- Graph $G = (V, E)$
 - **directed edges**: $e = (\text{source}, \text{destination})$
 - each edge and vertex **associated with a value** (user-defined type)
 - vertex and edge **values can be modified**
 - (structure modification also supported)




Terms: e is an out-edge of A, and in-edge of B.



Vertex-centric Programming

- “Think like a vertex”
- Popularized by the Pregel and GraphLab projects
 - Historically, systolic computation and the Connection Machine



MyFunc(vertex)
{ // modify neighborhood
}

The Main Challenge of Disk-based Graph Computation:

Random Access

Random Access Problem



- Symmetrized adjacency file with values,

vertex	in-neighbors	out-neighbors
5	3:2.3, 19: 1.3, 49: 0.65,...	781: 2.3, 881: 4.2...
....		
19	3: 1.4, 9: 12.1, ...	5: 1.3, 28: 2.2, ...

synchronize

Random
write

- ... or with millions of random accesses / second would be needed. Even for SSD, this is too much.


vertex	in-neighbors	out-neighbors
5	3: 881, 9: 12.1, ...	781: 2.3, 881: 4.2...
....		
19	3: 882, 9: 2872, ...	5: 1.3, 28: 2.2, ...

read

Random
read

Parallel Sliding Windows: Phases

- PSW processes the graph one **sub-graph** a time:




- In one **iteration**, the whole graph is processed.
 - And typically, next iteration is started.

1. Load
2. Compute
3. Write


PSW: Shards and Intervals

- Vertices are numbered from 1 to n
 - P intervals, each associated with a **shard** on disk.
 - **sub-graph** = interval of vertices




Example

Shard 1			Shard 2			Shard 3		
src	dst	value	src	dst	value	src	dst	value
1	2	0.3	1	3	0.4	2	5	0.6
3	2	0.2	2	3	0.3	3	5	0.9
4	1	1.4	3	4	0.8	4	6	1.2
5	1	0.5	5	3	0.2	5	5	0.3
6	2	0.6	6	4	1.9	6	6	1.1
	2	0.8						



(a) Execution interval (vertices 1-2)


Shard 1			Shard 2			Shard 3		
src	dst	value	src	dst	value	src	dst	value
1	2	0.273	1	3	0.364	2	5	0.545
3	2	0.22	2	3	0.273	3	5	0.9
4	1	1.54	3	4	0.8	4	6	1.2
5	1	0.55	5	3	0.2	5	5	0.3
6	2	0.66	6	4	1.9	6	6	1.1
	2	0.88						



(c) Execution interval (vertices 3-4)

(d) Execution interval (vertices 3-4)

(b) Execution interval (vertices 1-2)




PSW: Layout

1. Load

2. Compute

3. Write

Shard: in-edges for interval of vertices; sorted by source-id



Shards small enough to fit in memory; balance size of shards


PSW: Loading Sub-graph

Load subgraph for vertices 1..100

1. Load

2. Compute

3. Write



Load all in-edges in memory

What about out-edges?
Arranged in sequence in other shards

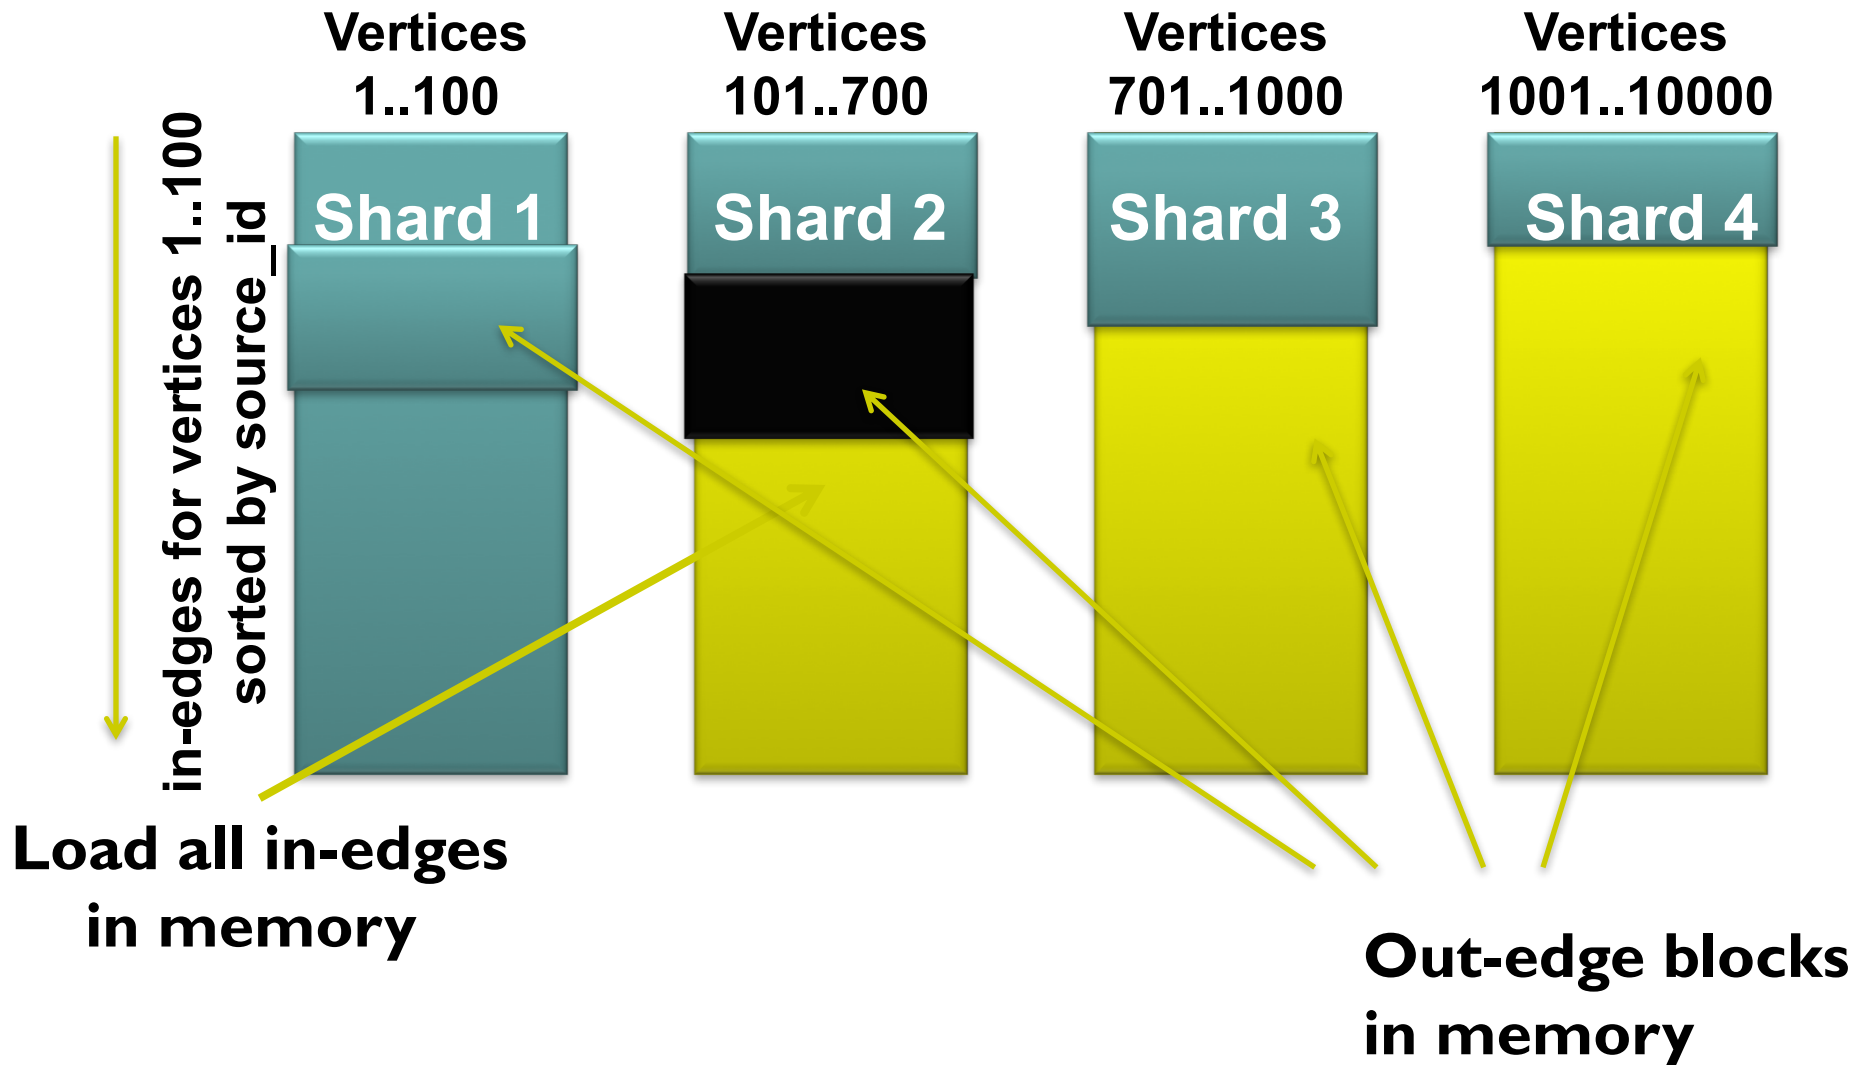
PSW: Loading Sub-graph

1. Load

2. Compute

3. Write

Load subgraph for vertices 101..700




PSW Load-Phase

1. Load

2. Compute

3. Write

Only P large reads for each interval.
 P^2 reads on one full pass.



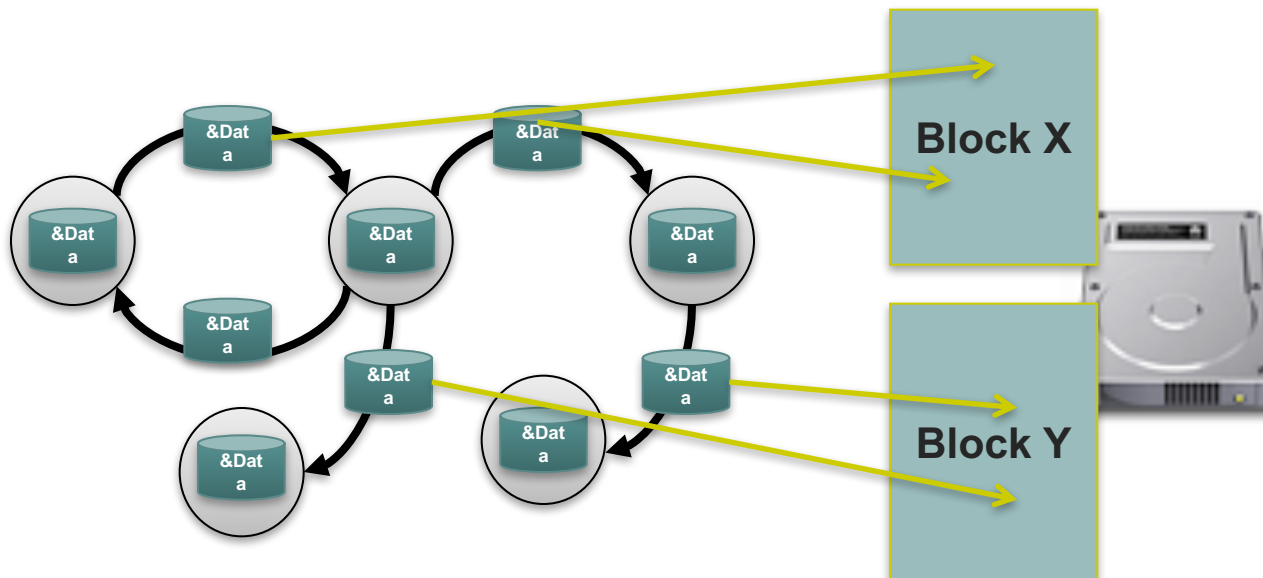
1. Load

2. Compute

3. Write

PSW: Execute updates

- Update-function is executed on interval's vertices
- Edges have pointers to the loaded data blocks
 - Changes take effect immediately → asynchronous.



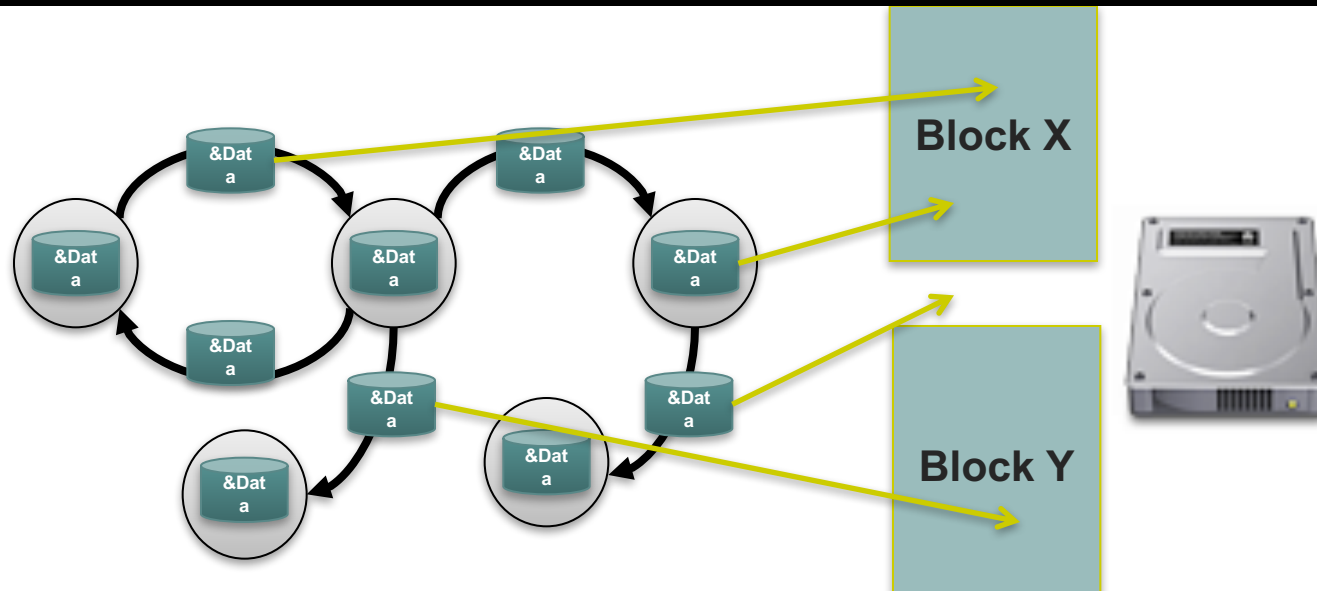
Deterministic scheduling prevents races between neighboring vertices.

1. Load
2. Compute
3. Write

PSW: Commit to Disk

- In write phase, the blocks are written *back* to disk
 - Next load-phase sees the preceding writes → asynchronous.

In total:
 P^2 reads and writes / full pass on the graph.
→ Performs well on *both* SSD and hard drive.



Programming

virtual void before_iteration(int iteration, graphchi_context &gcontext)

virtual void after_iteration(int iteration, graphchi_context &gcontext)

virtual bool repeat_updates(graphchi_context &gcontext)

virtual void before_exec_interval(vid_t window_st, vid_t window_en, graphchi_context &gcontext)

virtual void after_exec_interval(vid_t window_st, vid_t window_en, graphchi_context &gcontext)

virtual void update(vertex_t &v, graphchi_context &gcontext)=0

Evolving Graphs

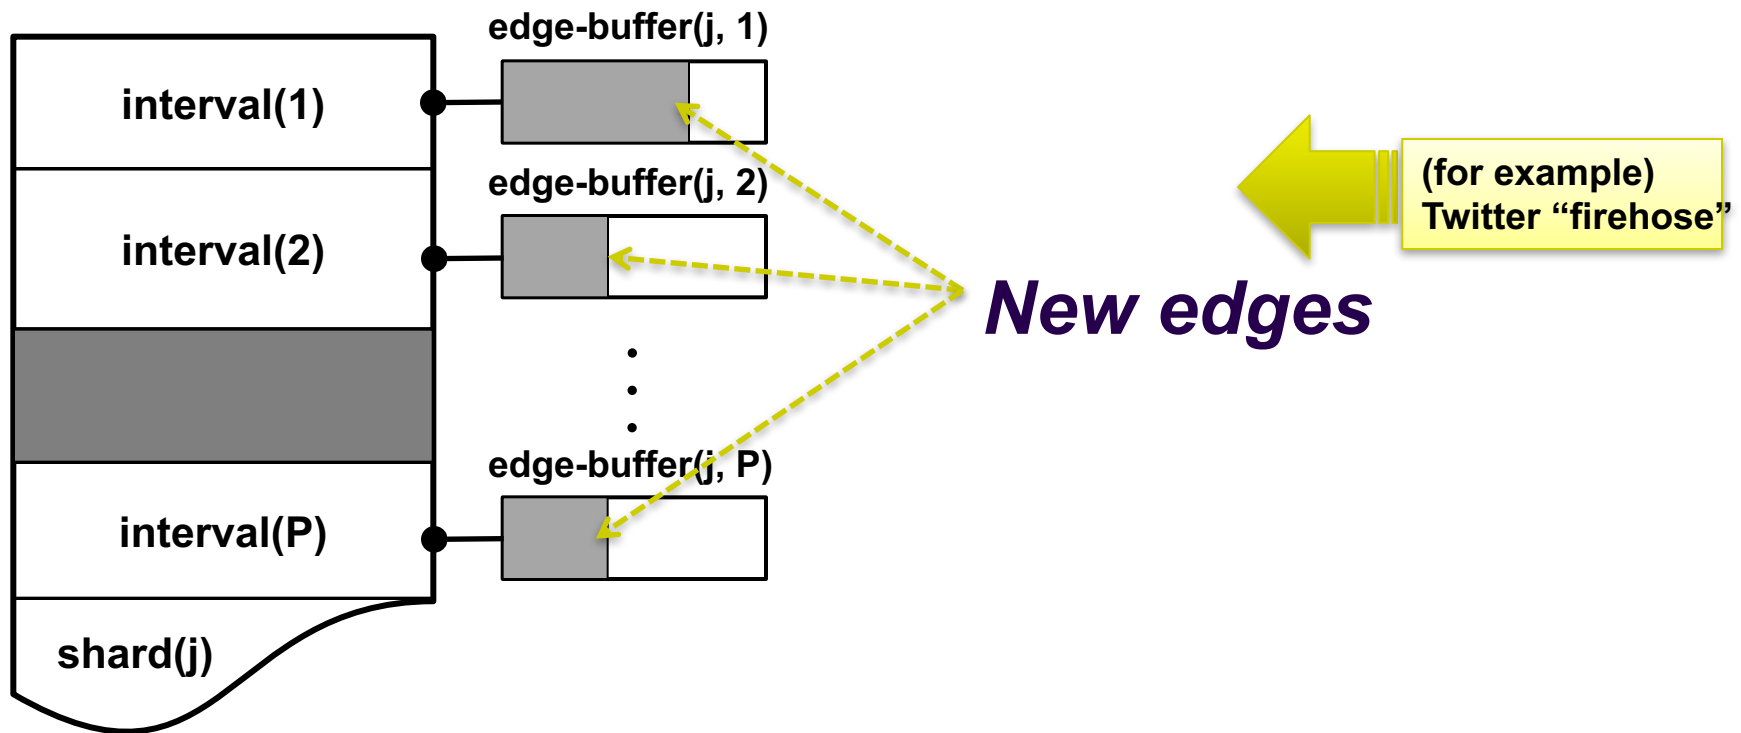
Graphs whose structure changes
over time

Evolving Graphs: Introduction

- Most interesting networks grow continuously:
 - New connections made, some ‘unfriended’.
- Desired functionality:
 - Ability to add and remove edges in streaming fashion;
 - ... while continuing computation.
- Related work:
 - *Kineograph* (EuroSys ‘12), distributed system for computation on a changing graph.

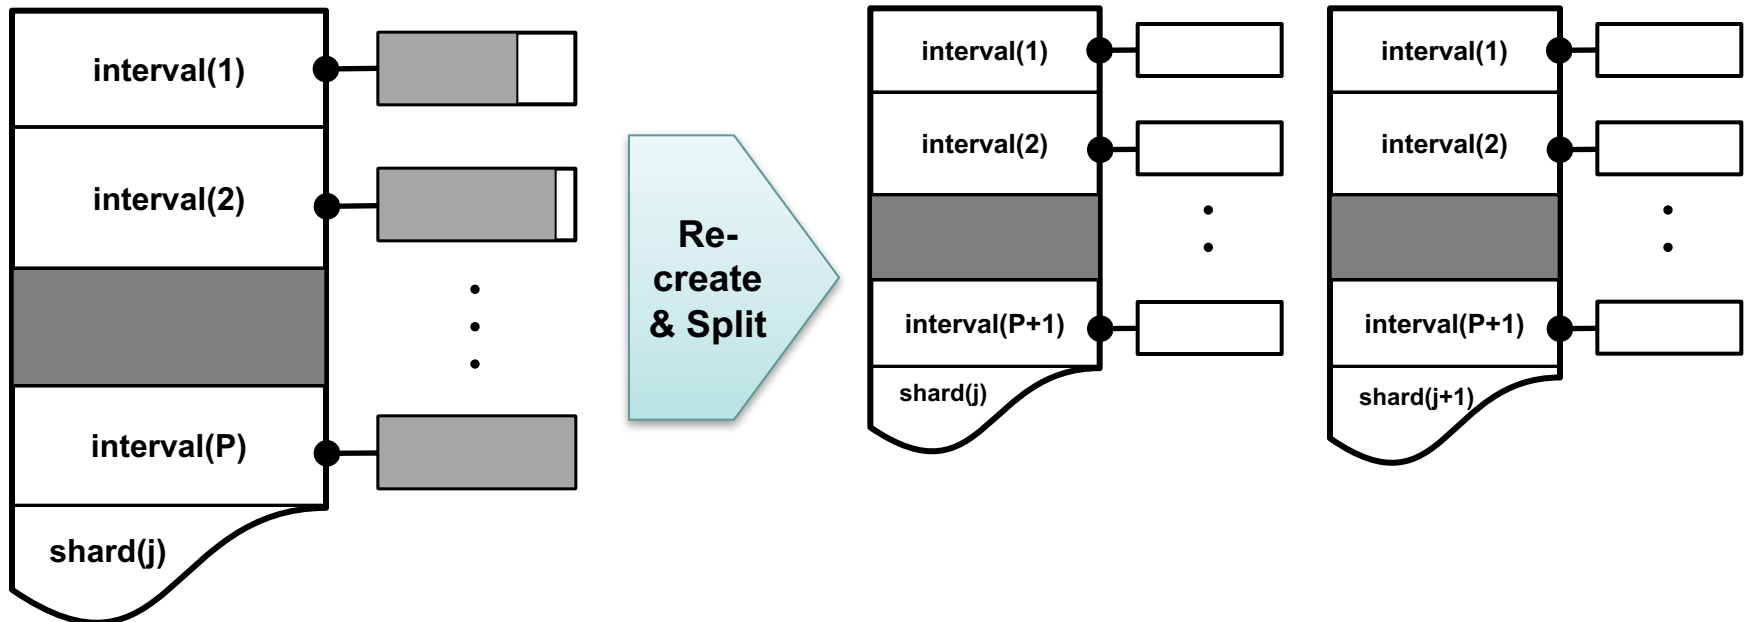
PSW and Evolving Graphs

- Adding edges
 - Each (shard, interval) has an associated edge-buffer.
- Removing edges: Edge flagged as “removed”.



Recreating Shards on Disk

- When buffers fill up, shards are **recreated** on disk
 - Too big shards are **split**.
- During recreation, deleted edges are permanently removed.



Distributed Graph Systems

- Pregel
- GraphLab
- GraphX